


What is a Shell?

- ▶ A program!
- ▶ A mechanism to interact with the computer / OS directly
- ▶ The CLI is just one type of shell



Shell Psuedocode

```
while(1){
    printf("person@machine$ ")

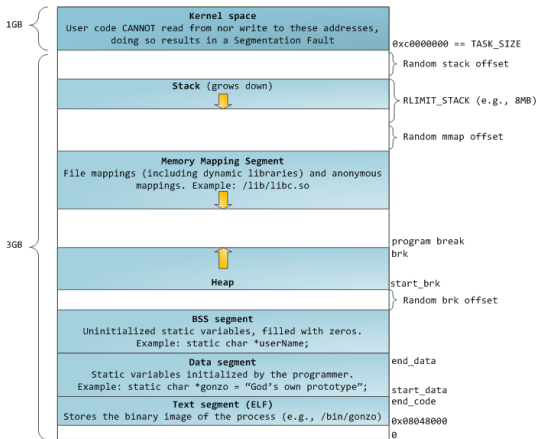
    // read_command();
    // cmd = command parsed
    // params = parameters parsed

    if(is_built_in(cmd))
    {
        do_built_in(cmd);
    }

    else
    {
        if(fork() != 0){
            waitpid(...?); // Parent
        } else {
            execve(cmd, params, 0); // Child
        }
    }
}
```

What Is A System Call?

- ▶ Functions provided by OS/kernel
- ▶ Run in kernel space
- ▶ Do important work for processes
- ▶ An interface to the machine



read Man Page

```
ed : man — Konsole
File Edit View Bookmarks Settings Help
READ(2)                                Linux Programmer's Manual                                READ(2)

NAME
    read - read from a file descriptor

SYNOPSIS
    #include <unistd.h>

    ssize_t read(int fd, void *buf, size_t count);

DESCRIPTION
    read() attempts to read up to count bytes from file descriptor fd into the buffer starting at buf.

    On files that support seeking, the read operation commences at the file offset, and the file offset is incremented by the number of bytes read. If the file offset is at or past the end of file, no bytes are read, and read() returns zero.

    If count is zero, read() may detect the errors described below. In the absence of any errors, or if read() does not check for errors, a read() with a count of 0 returns zero and has no other effects.

    According to POSIX.1, if count is greater than SSIZE_MAX, the result is implementation-
    Manual page read(2) line 1 (press h for help or q to quit)
```

Example System Calls

- ▶ chown
- ▶ chdir
- ▶ open
- ▶ read
- ▶ write
- ▶ close

- ▶ fork
- ▶ execve
- ▶ exit
- ▶ ioctl
- ▶ lseek
- ▶ mkdir

- ▶ mmap
- ▶ pipe
- ▶ reboot
- ▶ shutdown
- ▶ ...

Man Pages

```
File Edit View Bookmarks Settings Help
FORK(2) Linux Programmer's Manual FORK(2)

NAME
fork - create a child process

SYNOPSIS
#include <sys/types.h>
#include <unistd.h>

pid_t fork(void);

DESCRIPTION
fork() creates a new process by duplicating the calling process. The new process is referred to as the child process. The calling process is referred to as the parent process.

The child process and the parent process run in separate memory spaces. At the time of fork() both memory spaces have the same content. Memory writes, file mappings (mmap(2)), and unmappings (munmap(2)) performed by one of the processes do not affect the other.

The child process is an exact duplicate of the parent process except for the following points:

Manual page fork(2) line 1 (press h for help or q to quit)

ed:man

File Edit View Bookmarks Settings Help
EXECVE(2) Linux Programmer's Manual EXECVE(2)

NAME
execve - execute program

SYNOPSIS
#include <unistd.h>

int execve(const char *filename, char *const argv[],
           char *const envp[]);

DESCRIPTION
execve() executes the program pointed to by filename. filename must be either a binary executable, or a script starting with a line of the form:

#! interpreter [optional-arg]

For details of the latter case, see "Interpreter scripts" below.

argv is an array of argument strings passed to the new program. By convention, the first of these strings (i.e., argv[0]) should contain the filename associated with the file being executed. envp is an array of strings, conventionally of the form key=value, which are passed as environment to the new program. The argv and envp arrays must each include

Manual page execve(2) line 1 (press h for help or q to quit)

ed:man
```

Fork Man Page

```
File Edit View Bookmarks Settings Help
FORK(2) Linux Programmer's Manual FORK(2)

NAME
    fork - create a child process

SYNOPSIS
    #include <sys/types.h>
    #include <unistd.h>

    pid_t fork(void);

DESCRIPTION
    fork() creates a new process by duplicating the calling process. The new process is
    referred to as the child process. The calling process is referred to as the parent
    process.

    The child process and the parent process run in separate memory spaces. At the time of
    fork() both memory spaces have the same content. Memory writes, file mappings (mmap(2)),
    and unmappings (munmap(2)) performed by one of the processes do not affect the other.

    The child process is an exact duplicate of the parent process except for the following
    points:

Manual page fork(2) line 1 (press h for help or q to quit)

ed : man
```

Execve Man Page

```
File Edit View Bookmarks Settings Help
EXECVE(2)                                Linux Programmer's Manual                                EXECVE(2)

NAME
    execve - execute program

SYNOPSIS
    #include <unistd.h>

    int execve(const char *filename, char *const argv[],
               char *const envp[]);

DESCRIPTION
    execve() executes the program pointed to by filename. filename must be either a binary
    executable, or a script starting with a line of the form:

        #! interpreter [optional-arg]

    For details of the latter case, see "Interpreter scripts" below.

    argv is an array of argument strings passed to the new program. By convention, the first
    of these strings (i.e., argv[0]) should contain the filename associated with the file
    being executed. envp is an array of strings, conventionally of the form key=value, which
    are passed as environment to the new program. The argv and envp arrays must each include
    Manual page execve(2) line 1 (press h for help or q to quit)
```

ed : man

Fork() Diagram

